

# Weekly Summary Template

Brady Miller

## Table of contents

Tuesday, April 11 . . . . .	1
Tuesday April 11 . . . . .	1
Thursday, April 13 . . . . .	17

## Tuesday, April 11

### ! TIL

Include a *very brief* summary of what you learnt in this class here.

Today, I learnt the following concepts in class:

1. Application of neural networks to breast cancer & titanic datasets
2. What are DataLoaders and its strengths
3. How to use DataLoaders to determine the value of a handwritten number

## Tuesday April 11

```
packages <- c(
  # Old packages
  "ISLR2",
  "dplyr",
  "tidyr",
  "readr",
  "purrr",
  "repr",
  "tidyverse",
  "kableExtra",
```

```
  "IRdisplay",
  # NEW
  "torch",
  "torchvision",
  "luz"
)

# renv::install(packages)
sapply(packages, require, character.only=TRUE)
```

Loading required package: ISLR2

Loading required package: dplyr

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

Loading required package: tidyr

Loading required package: readr

Loading required package: purrr

Loading required package: repr

Loading required package: tidyverse

Warning in library(package, lib.loc = lib.loc, character.only = TRUE,  
logical.return = TRUE, : there is no package called 'tidyverse'

Loading required package: kableExtra

Attaching package: 'kableExtra'

The following object is masked from 'package:dplyr':

group\_rows

Loading required package: IRdisplay

Loading required package: torch

Loading required package: torchvision

Loading required package: luz

ISLR2	dplyr	tidyr	readr	purrr	repr
TRUE	TRUE	TRUE	TRUE	TRUE	TRUE
tidyverse	kableExtra	IRdisplay	torch	torchvision	luz
FALSE	TRUE	TRUE	TRUE	TRUE	TRUE

## Application of neural networks to breast cancer & Titanic datasets

```
url <- "https://web.stanford.edu/class/archive/cs/cs109/cs109.1166/stuff/titanic.csv"

df <- read_csv(url) %>%
  mutate_if(\(x) is.character(x), as.factor) %>%
  mutate(y = Survived) %>%
  select(-c(Name, Survived)) %>%
  (\(x) {
    names(x) <- tolower(names(x))
    x
  })
```

Rows: 887 Columns: 8

-- Column specification -----  
Delimiter: ","

chr (2): Name, Sex

dbl (6): Survived, Pclass, Age, Siblings/Spouses Aboard, Parents/Children Ab...

i Use `spec()` to retrieve the full column specification for this data.

i Specify the column types or set `show\_col\_types = FALSE` to quiet this message.

```

# url <- "https://archive.ics.uci.edu/ml/machine-learning-databases/breast-cancer-wisconsin

# col_names <- c("id", "diagnosis", paste0("feat", 1:30))

# df <- read_csv(
#   url, col_names, col_types = cols()
# ) %>%
#   select(-id) %>%
#   mutate(y = ifelse(diagnosis == "M", 1, 0)) %>%
#   select(-diagnosis)

# df %>% head

```

```

k <- 5

test_ind <- sample(
  1:nrow(df),
  floor(nrow(df) / k),
  replace=FALSE
)

df_train <- df[-test_ind, ]
df_test  <- df[test_ind, ]

nrow(df_train) + nrow(df_test) == nrow(df)

```

[1] TRUE

```

fit_glm <- glm(
  y ~ .,
  df_train %>% mutate_at("y", factor),
  family = binomial()
)

glm_test <- predict(
  fit_glm,
  df_test,
  output = "response"
)

```

```
glm_preds <- ifelse(glm_test > 0.5, 1, 0)
table(glm_preds, df_test$y)
```

```
glm_preds  0  1
          0 98 26
          1  9 44
```

```
NNet <- nn_module(
  initialize = function(p, q1, q2, q3) {
    self$hidden1 <- nn_linear(p, q1)
    self$hidden2 <- nn_linear(q1, q2)
    self$hidden3 <- nn_linear(q2, q3)
    self$output <- nn_linear(q3, 1)
    self$activation <- nn_relu()
    self$sigmoid <- nn_sigmoid()
  },

  forward = function(x) {
    x %>%
      self$hidden1() %>% self$activation() %>%
      self$hidden2() %>% self$activation() %>%
      self$hidden3() %>% self$activation() %>%
      self$output() %>% self$sigmoid()
  }
)
```

```
# fitting a model without an intercept - when its zero, everything else is zero
M <- model.matrix(y ~ 0 + ., data = df_train)
# model.matrix(y ~ ., data = df_train) [ , -1]
```

```
fit_nn <- NNet %>%
  #
  # Setup the model
  #
  setup(
    loss = nn_bce_loss(),
    optimizer = optim_adam,
    metrics = list(
      luz_metric_accuracy()
    )
  ) %>%
```

```

#
# Set the hyperparameters
#
set_hparams(p=ncol(M), q1=256, q2=128, q3=64) %>%
set_opt_hparams(lr=0.005) %>%
#
# Fit the model
#
fit(
  data = list(
    model.matrix(y ~ 0 + ., data = df_train),
    df_train %>% select(y) %>% as.matrix
  ),
  valid_data = list(
    model.matrix(y ~ 0 + ., data = df_test),
    df_test %>% select(y) %>% as.matrix
  ),
  epochs = 50,
  verbose = TRUE
)

```

```

Epoch 1/50
Train metrics: Loss: 0.7377 - Acc: 12.1127
Valid metrics: Loss: 0.6321 - Acc: 12.2316
Epoch 2/50
Train metrics: Loss: 0.641 - Acc: 12.1859
Valid metrics: Loss: 0.5933 - Acc: 12.2316
Epoch 3/50
Train metrics: Loss: 0.6094 - Acc: 12.1859
Valid metrics: Loss: 0.5666 - Acc: 12.2316
Epoch 4/50
Train metrics: Loss: 0.6024 - Acc: 12.1859
Valid metrics: Loss: 0.565 - Acc: 12.2316
Epoch 5/50
Train metrics: Loss: 0.5987 - Acc: 12.1493
Valid metrics: Loss: 0.5409 - Acc: 12.2316
Epoch 6/50
Train metrics: Loss: 0.5818 - Acc: 12.2592
Valid metrics: Loss: 0.525 - Acc: 12.2316
Epoch 7/50
Train metrics: Loss: 0.5505 - Acc: 12.1493
Valid metrics: Loss: 0.4909 - Acc: 12.2316

```


Epoch 8/50  
Train metrics: Loss: 0.5624 - Acc: 12.1493  
Valid metrics: Loss: 0.5084 - Acc: 12.2316  
Epoch 9/50  
Train metrics: Loss: 0.5272 - Acc: 12.1493  
Valid metrics: Loss: 0.4551 - Acc: 12.2316  
Epoch 10/50  
Train metrics: Loss: 0.5086 - Acc: 12.2225  
Valid metrics: Loss: 0.5302 - Acc: 12.2316  
Epoch 11/50  
Train metrics: Loss: 0.5148 - Acc: 12.1493  
Valid metrics: Loss: 0.4628 - Acc: 12.2316  
Epoch 12/50  
Train metrics: Loss: 0.4761 - Acc: 12.1859  
Valid metrics: Loss: 0.445 - Acc: 12.2316  
Epoch 13/50  
Train metrics: Loss: 0.4583 - Acc: 12.1859  
Valid metrics: Loss: 0.5348 - Acc: 12.2316  
Epoch 14/50  
Train metrics: Loss: 0.4715 - Acc: 12.1493  
Valid metrics: Loss: 0.4695 - Acc: 12.2316  
Epoch 15/50  
Train metrics: Loss: 0.4781 - Acc: 12.1859  
Valid metrics: Loss: 0.4257 - Acc: 12.2316  
Epoch 16/50  
Train metrics: Loss: 0.4439 - Acc: 12.2225  
Valid metrics: Loss: 0.4369 - Acc: 12.2316  
Epoch 17/50  
Train metrics: Loss: 0.4566 - Acc: 12.1493  
Valid metrics: Loss: 0.4609 - Acc: 12.2316  
Epoch 18/50  
Train metrics: Loss: 0.4246 - Acc: 12.1859  
Valid metrics: Loss: 0.4646 - Acc: 12.2316  
Epoch 19/50  
Train metrics: Loss: 0.4278 - Acc: 12.1859  
Valid metrics: Loss: 0.4317 - Acc: 12.2316  
Epoch 20/50  
Train metrics: Loss: 0.412 - Acc: 12.1127  
Valid metrics: Loss: 0.453 - Acc: 12.2316  
Epoch 21/50  
Train metrics: Loss: 0.4501 - Acc: 12.1859  
Valid metrics: Loss: 0.4772 - Acc: 12.2316  
Epoch 22/50

Train metrics: Loss: 0.4231 - Acc: 12.2225  
Valid metrics: Loss: 0.4891 - Acc: 12.2316  
Epoch 23/50  
Train metrics: Loss: 0.4306 - Acc: 12.1859  
Valid metrics: Loss: 0.4389 - Acc: 12.2316  
Epoch 24/50  
Train metrics: Loss: 0.4305 - Acc: 12.1493  
Valid metrics: Loss: 0.4127 - Acc: 12.2316  
Epoch 25/50  
Train metrics: Loss: 0.4267 - Acc: 12.2225  
Valid metrics: Loss: 0.4122 - Acc: 12.2316  
Epoch 26/50  
Train metrics: Loss: 0.4159 - Acc: 12.2225  
Valid metrics: Loss: 0.4006 - Acc: 12.2316  
Epoch 27/50  
Train metrics: Loss: 0.4064 - Acc: 12.1859  
Valid metrics: Loss: 0.4176 - Acc: 12.2316  
Epoch 28/50  
Train metrics: Loss: 0.4097 - Acc: 12.1859  
Valid metrics: Loss: 0.4163 - Acc: 12.2316  
Epoch 29/50  
Train metrics: Loss: 0.4049 - Acc: 12.2225  
Valid metrics: Loss: 0.4363 - Acc: 12.2316  
Epoch 30/50  
Train metrics: Loss: 0.4314 - Acc: 12.1859  
Valid metrics: Loss: 0.4184 - Acc: 12.2316  
Epoch 31/50  
Train metrics: Loss: 0.4244 - Acc: 12.1859  
Valid metrics: Loss: 0.4377 - Acc: 12.2316  
Epoch 32/50  
Train metrics: Loss: 0.4258 - Acc: 12.1859  
Valid metrics: Loss: 0.4417 - Acc: 12.2316  
Epoch 33/50  
Train metrics: Loss: 0.397 - Acc: 12.1493  
Valid metrics: Loss: 0.4213 - Acc: 12.2316  
Epoch 34/50  
Train metrics: Loss: 0.3922 - Acc: 12.2225  
Valid metrics: Loss: 0.4252 - Acc: 12.2316  
Epoch 35/50  
Train metrics: Loss: 0.3924 - Acc: 12.2225  
Valid metrics: Loss: 0.4338 - Acc: 12.2316  
Epoch 36/50  
Train metrics: Loss: 0.3975 - Acc: 12.1859



Valid metrics: Loss: 0.4246 - Acc: 12.2316  
Epoch 37/50  
Train metrics: Loss: 0.3994 - Acc: 12.2592  
Valid metrics: Loss: 0.4518 - Acc: 12.2316  
Epoch 38/50  
Train metrics: Loss: 0.4062 - Acc: 12.1493  
Valid metrics: Loss: 0.4811 - Acc: 12.2316  
Epoch 39/50  
Train metrics: Loss: 0.403 - Acc: 12.1127  
Valid metrics: Loss: 0.4329 - Acc: 12.2316  
Epoch 40/50  
Train metrics: Loss: 0.388 - Acc: 12.1859  
Valid metrics: Loss: 0.4559 - Acc: 12.2316  
Epoch 41/50  
Train metrics: Loss: 0.4037 - Acc: 12.1859  
Valid metrics: Loss: 0.4479 - Acc: 12.2316  
Epoch 42/50  
Train metrics: Loss: 0.4046 - Acc: 12.2225  
Valid metrics: Loss: 0.4666 - Acc: 12.2316  
Epoch 43/50  
Train metrics: Loss: 0.3939 - Acc: 12.2592  
Valid metrics: Loss: 0.4423 - Acc: 12.2316  
Epoch 44/50  
Train metrics: Loss: 0.403 - Acc: 12.1493  
Valid metrics: Loss: 0.4708 - Acc: 12.2316  
Epoch 45/50  
Train metrics: Loss: 0.4048 - Acc: 12.1493  
Valid metrics: Loss: 0.4531 - Acc: 12.2316  
Epoch 46/50  
Train metrics: Loss: 0.3916 - Acc: 12.1859  
Valid metrics: Loss: 0.4634 - Acc: 12.2316  
Epoch 47/50  
Train metrics: Loss: 0.3926 - Acc: 12.1493  
Valid metrics: Loss: 0.4533 - Acc: 12.2316  
Epoch 48/50  
Train metrics: Loss: 0.4003 - Acc: 12.2225  
Valid metrics: Loss: 0.4536 - Acc: 12.2316  
Epoch 49/50  
Train metrics: Loss: 0.4079 - Acc: 12.1859  
Valid metrics: Loss: 0.4405 - Acc: 12.2316  
Epoch 50/50  
Train metrics: Loss: 0.3881 - Acc: 12.1859  
Valid metrics: Loss: 0.4579 - Acc: 12.2316

```
plot(fit_nn)
```



```
nn_test <- predict(
  fit_nn,
  model.matrix(y ~ . - 1, data = df_test)
)
# nn_test
nn_preds <- ifelse(nn_test > 0.5, 1, 0)

table(nn_preds, df_test$y)
```

```
nn_preds 0 1
0 98 29
1 9 41
```

```
mean(nn_preds == df_test$y)
```

```
[1] 0.7853107
```

```
table(glm_preds, df_test$y)
```

```
glm_preds  0  1  
          0 98 26  
          1  9 44
```

```
mean(glm_preds == df_test$y)
```

```
[1] 0.8022599
```

## DataLoaders

- key component in the ML pipeline
- handle loading and preprocessing of data in efficient way for training and evaluating models
- make it easy to work with large datasets (ex. dataset with trillions of bits)
  - Loading data in smaller batches called chunks
- Advantages
  1. Efficiency memory management
  2. Parallelism
  3. Preprocessing
  4. Flexibility
  5. Standardization
- involves partitioning data into smaller chunks (batches)
- inside one large epoch of gradient descent, it'll break up the data into batches and mini gradient descents (does gradient descent on every batch)
- IS DIFFERENT FROM CROSS VALIDATION
- training data is split up into smaller chunks - data is being cycled through during gradient descent
- if batches are fixed statically (first 100 are batch 1, next 100 are batch 2), this is going to introduce some bias into the gradient descent process, so you want to shuffle the data between every epoch
- DataLoader takes care of these issues

```
dir <- "./mnist"
```

```
train_ds <- mnist_dataset(
```

```

    root = dir,
    train = TRUE,
    download = TRUE,
    transform = transform
  )
test_ds <- mnist_dataset(
  root = dir,
  train = FALSE,
  download = TRUE,
  transform = transform
)

```

```
typeof(train_ds)
```

```
[1] "environment"
```

```
length(train_ds)
```

```
[1] 60000
```

```

# 42,000nd observation is a 28x28 matrix
# one row is a 28x28 matrix
# is an image where a value specifies the intensity of the pixels in the 28x28 image
train_ds$data[42000, ]

```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]	[,11]	[,12]	[,13]
[1,]	0	0	0	0	0	0	0	0	0	0	0	0	0
[2,]	0	0	0	0	0	0	0	0	0	0	0	0	0
[3,]	0	0	0	0	0	0	0	0	0	0	0	0	0
[4,]	0	0	0	0	0	0	0	0	0	0	0	0	0
[5,]	0	0	0	0	0	0	0	0	0	0	0	0	0
[6,]	0	0	0	0	0	0	0	0	0	0	0	0	0
[7,]	0	0	0	0	0	0	0	0	0	0	0	6	37
[8,]	0	0	0	0	0	0	0	0	0	0	125	160	252
[9,]	0	0	0	0	0	0	0	0	1	109	232	252	252
[10,]	0	0	0	0	0	0	0	0	125	252	252	252	252
[11,]	0	0	0	0	0	0	0	0	62	189	211	252	252
[12,]	0	0	0	0	0	0	0	21	206	252	190	252	168

[13,]	0	0	0	0	0	0	73	253	253	253	253	217	0
[14,]	0	0	0	0	0	0	115	252	252	252	148	30	0
[15,]	0	0	0	0	0	0	217	252	252	252	35	0	0
[16,]	0	0	0	0	0	0	217	252	252	252	35	0	0
[17,]	0	0	0	0	0	110	233	253	253	144	0	79	109
[18,]	0	0	0	0	0	253	252	252	252	237	217	242	252
[19,]	0	0	0	0	0	253	252	252	252	252	252	252	252
[20,]	0	0	0	0	0	170	252	252	252	252	252	252	252
[21,]	0	0	0	0	0	0	218	253	253	253	253	253	253
[22,]	0	0	0	0	0	0	72	231	252	252	252	252	252
[23,]	0	0	0	0	0	0	0	52	71	71	71	71	71
[24,]	0	0	0	0	0	0	0	0	0	0	0	0	0
[25,]	0	0	0	0	0	0	0	0	0	0	0	0	0
[26,]	0	0	0	0	0	0	0	0	0	0	0	0	0
[27,]	0	0	0	0	0	0	0	0	0	0	0	0	0
[28,]	0	0	0	0	0	0	0	0	0	0	0	0	0
[,14] [,15] [,16] [,17] [,18] [,19] [,20] [,21] [,22] [,23] [,24] [,25]													
[1,]	0	0	0	0	0	0	0	0	0	0	0	0	0
[2,]	0	0	0	0	0	0	0	0	0	0	0	0	0
[3,]	0	0	0	0	0	0	0	0	0	0	0	0	0
[4,]	0	0	0	0	0	0	0	0	0	0	0	0	0
[5,]	0	0	0	0	0	0	0	0	0	0	0	0	0
[6,]	0	0	0	0	0	0	0	42	218	134	186	0	0
[7,]	182	98	51	0	0	0	27	221	253	252	221	16	0
[8,]	253	252	175	144	0	0	16	190	253	252	252	108	0
[9,]	253	252	252	252	0	0	0	0	109	252	236	62	0
[10,]	253	252	200	179	0	0	0	0	109	252	215	42	0
[11,]	237	91	20	0	0	0	0	21	212	252	241	221	0
[12,]	62	0	0	0	0	0	21	206	253	252	252	252	0
[13,]	0	0	0	0	0	32	212	253	255	253	253	108	0
[14,]	0	0	0	0	0	115	252	252	253	252	220	15	0
[15,]	0	0	27	120	182	242	252	252	253	252	112	0	0
[16,]	0	125	221	252	253	252	252	252	253	128	31	0	0
[17,]	255	253	253	253	255	253	253	253	208	20	0	0	0
[18,]	253	252	252	252	253	252	252	210	20	0	0	0	0
[19,]	253	252	252	252	217	215	112	31	0	0	0	0	0
[20,]	253	252	252	252	0	0	0	0	0	0	0	0	0
[21,]	255	253	175	62	0	0	0	0	0	0	0	0	0
[22,]	119	35	10	0	0	0	0	0	0	0	0	0	0
[23,]	0	0	0	0	0	0	0	0	0	0	0	0	0
[24,]	0	0	0	0	0	0	0	0	0	0	0	0	0
[25,]	0	0	0	0	0	0	0	0	0	0	0	0	0
[26,]	0	0	0	0	0	0	0	0	0	0	0	0	0

[27,]	0	0	0	0	0	0	0	0	0	0	0	0
[28,]	0	0	0	0	0	0	0	0	0	0	0	0
	[,26]	[,27]	[,28]									
[1,]	0	0	0									
[2,]	0	0	0									
[3,]	0	0	0									
[4,]	0	0	0									
[5,]	0	0	0									
[6,]	0	0	0									
[7,]	0	0	0									
[8,]	0	0	0									
[9,]	0	0	0									
[10,]	0	0	0									
[11,]	0	0	0									
[12,]	0	0	0									
[13,]	0	0	0									
[14,]	0	0	0									
[15,]	0	0	0									
[16,]	0	0	0									
[17,]	0	0	0									
[18,]	0	0	0									
[19,]	0	0	0									
[20,]	0	0	0									
[21,]	0	0	0									
[22,]	0	0	0									
[23,]	0	0	0									
[24,]	0	0	0									
[25,]	0	0	0									
[26,]	0	0	0									
[27,]	0	0	0									
[28,]	0	0	0									


### Using the data/DataLoaders to determine the value of a handwritten number

```
# collection of {x_i, y_i} where i = 1, ranging to 10000
# x_i is 28x28 image of handwritten image
# every y_i {0,1,...,9}
options(repr.plot.width=10, repr.plot.height=10)

i <- sample(1:length(train_ds), 1)
x <- train_ds$data[i, ] %>% t
```

```
image(x[1:28, 28:1], useRaster=TRUE, axes=FALSE, col=gray.colors(1000), main = train_ds$ta
```

4



- data frame with each row being a single image observation and its actual value
- $p$  (# cols) = # of pixels in  $\text{width}^2$  so you can have column for each pixel

```
# splitting up batches into size 128
# shuffling allows to negate any bias by shuffling observations for each epoch
train_dl <- dataloader(train_ds, batch_size = 1024, shuffle = TRUE)
test_dl <- dataloader(test_ds, batch_size = 1024)
```

```
NNet_10 <- nn_module(
  initialize = function(p, q1, q2, q3, o) {
    self$hidden1 <- nn_linear(p, q1)
    self$hidden2 <- nn_linear(q1, q2)
    self$hidden3 <- nn_linear(q2, q3)
    self$OUTPUT <- nn_linear(q3, o)
    self$activation <- nn_relu()
  },
  forward = function(x) {
    x %>%
      self$hidden1() %>%
```

```

        self$activation() %>%
        self$hidden2() %>%
        self$activation() %>%
        self$hidden3() %>%
        self$activation() %>%
        self$OUTPUT()
    }
)

```

Unable to get this chunk to work as it gives error that says it can't convert argument: input

```

fit_nn <- NNet_10 %>%
  #
  # Setup the model
  #
  setup(
    loss = nn_cross_entropy_loss(),
    optimizer = optim_adam,
    metrics = list(
      luz_metric_accuracy()
    )
  ) %>%
  #
  # Set the hyperparameters
  #
  set_hparams(p=28*28, q1=256, q2=128, q3=64, o=10) %>%
  #
  # Fit the model
  #
  fit(
    epochs = 10,
    data = train_dl,
    valid_data = test_dl,
    verbose=TRUE
  )

NN10_preds <- fit_nn %>%
  predict(test_ds) %>%
  torch_argmax(dim = 2) %>%
  as_array()

```



**Thursday, April 13**

! TIL

Include a *very brief* summary of what you learnt in this class here.

Today, I learnt the following concepts in class:

1. Item 1
2. Item 2
3. Item 3